Computer Vision-Based Cancer Detection Using CNN

Aditya Adhana

May 17, 2025

Abstract

This report details a deep learning methodology for cancer detection leveraging computer vision techniques. By training a convolutional neural network (CNN) on meticulously labeled medical image datasets, our objective is to automate the classification of images as either cancerous or non-cancerous. The promising results obtained underscore the potential of AI-assisted diagnostics within the medical field.

1. Introduction

The diagnosis of cancer frequently involves the thorough examination of medical images. Traditional approaches to this analysis are often time-intensive and heavily reliant on the expertise of human interpreters. Recent breakthroughs in deep learning, particularly the development and application of convolutional neural networks (CNNs), have led to significant advancements in image classification tasks. This project centers on the development of a CNN model specifically designed to categorize medical images as cancerous or non-cancerous, utilizing publicly accessible datasets for training and evaluation.

2. Dataset Description

The dataset employed in this study comprises labeled medical images, conveniently stored on Google Drive and accessed through the Google Colab environment. The image collection was systematically organized into distinct training and validation directories, with each directory further subdivided into folders corresponding to the specific cancer and non-cancer classes.

3. Methodology

3.1 Data Preprocessing

Prior to training, the medical images underwent a crucial preprocessing step involving rescaling. This was achieved using Keras's `ImageDataGenerator`, which normalized the pixel intensity values to a standardized range between $[0, 1]$ (inclusive). This normalization process is vital as it facilitates faster and more stable convergence of the model during the training phase.

3.2 Model Architecture

The architecture of the convolutional neural network (CNN) developed for this project incorporates the following key layers, arranged sequentially:

- **Convolutional layers:** These layers are fundamental to feature extraction from the input images. Each convolutional layer is followed by a ReLU (Rectified Linear Unit) activation function, introducing non-linearity to the model.
- **MaxPooling layers:** These layers perform spatial downsampling, reducing the dimensionality of the feature maps while retaining the most salient information. This helps in making the model more robust to variations in the input images.
- **Fully connected Dense layers:** Following the convolutional and pooling layers, fully connected layers are used to perform the final classification based on the learned features.
- **Output layer:** The final layer is a Dense layer with a sigmoid activation function. The sigmoid function is particularly well-suited for binary classification tasks, as it outputs a probability between 0 and 1, representing the likelihood of the image belonging to the cancerous class.

3.3 Compilation and Training

The developed CNN model was compiled using the RMSprop optimization algorithm, known for its effectiveness in training deep neural networks. The binary cross-entropy loss function was chosen as the metric to be minimized during training, as it is standard for binary classification problems. The model was trained over a defined number of epochs, with both the training and validation datasets being fed into the model through data generators. These generators efficiently provide batches of preprocessed images to the model during the training process.

4. Implementation Details

The entire project was implemented using the Python programming language, leveraging the powerful deep learning libraries TensorFlow and Keras. In addition to these core libraries, OpenCV (Open Source Computer Vision Library) was utilized for various image manipulation tasks, and Matplotlib was employed for the visualization of data and results.

5. Results and Discussion

The trained convolutional neural network (CNN) demonstrated a satisfactory level of accuracy when evaluated on the held-out validation dataset. This outcome highlights the considerable potential of such models in assisting medical professionals with diagnostic tasks. Furthermore, the model's predictions on a set of unseen test images showed correct classification for the majority of samples. However, the presence of some misclassifications underscores the necessity for ongoing refinement of the model architecture, optimization of training parameters, and the incorporation of a more diverse and representative dataset to enhance its generalization capabilities.

6. Conclusion

This project has successfully provided a proof-of-concept for a cancer detection system based on deep learning and computer vision techniques. The findings suggest that with continued development, meticulous refinement, and thorough clinical validation, such intelligent systems have the potential to evolve into invaluable tools for medical diagnostics, ultimately aiding in earlier and more accurate cancer detection.

References

- Chollet, F. (2017). *Deep Learning with Python*. Manning.
- Keras Documentation: https://keras.io/
- TensorFlow Documentation: https://www.tensorflow.org/